

Multisensor-Fusion for 3D Full-Body Human Motion Capture

Gerard Pons-Moll¹, Andreas Baak², Thomas Helten², Meinard Müller², Hans-Peter Seidel², Bodo Rosenhahn¹

¹Leibniz Universität Hannover, Germany

²Saarland University & MPI Informatik, Germany

{pons, rosenhahn}@tnt.uni-hannover.de

{abaak, thelten, meinard}@mpi-inf.mpg.de

Abstract

In this work, we present an approach to fuse video with orientation data obtained from extended inertial sensors to improve and stabilize full-body human motion capture. Even though video data is a strong cue for motion analysis, tracking artifacts occur frequently due to ambiguities in the images, rapid motions, occlusions or noise. As a complementary data source, inertial sensors allow for drift-free estimation of limb orientations even under fast motions. However, accurate position information cannot be obtained in continuous operation. Therefore, we propose a hybrid tracker that combines video with a small number of inertial units to compensate for the drawbacks of each sensor type: on the one hand, we obtain drift-free and accurate position information from video data and, on the other hand, we obtain accurate limb orientations and good performance under fast motions from inertial sensors. In several experiments we demonstrate the increased performance and stability of our human motion tracker.

1. Introduction

In this paper, we deal with the task of human pose tracking, also known as motion capturing (MoCap) [14]. A basic prerequisite for our system is a 3D model of the person and at least one calibrated camera view. The goal of MoCap is to obtain the 3D pose of the person, which is in general an ambiguous problem. Using additional a priori knowledge such as familiar pose configurations learned from motion capture data helps considerably to handle more difficult scenarios like partial occlusions, background clutter, or corrupted image data. There are several ways to employ such a priori knowledge to human tracking. One option is to learn the space of plausible human poses and motions [2, 4, 12, 13, 21, 19]. Another option is to learn a direct mapping from image features to the pose space [1, 9, 19, 25]. To constrain the high dimensional space of kinematic models, a major theme of recent research on human tracking has been dealing with dimensionality reduction [27, 28]. Here, the idea is that a typical motion pat-




Figure 1: Tracking result for two selected frames. (a) Video-based tracker. (b) Our proposed hybrid tracker.

tern like walking should be a rather simple trajectory in a lower dimensional manifold. Therefore, prior distributions are learned in this lower dimensional space. Such methods are believed to generalize well with only little training data. Inspired by the same ideas of dimensionality reduction, physical and illumination models have been recently proposed to constrain and to represent human motion in a more realistic way [6, 3, 11, 23]. A current trend of research tries to estimate shape deformations from images besides the body pose by either directly deforming the mesh geometry [7] or by a combination of skeleton-based pose estimation with surface deformation [10].

Recently, inertial sensors (e.g. gyroscopes and accelerometers) have become popular for human motion analysis. Often, sensors are used for medical applications, see, e.g., [8] where accelerometer and gyroscope data is fused. However, their application concentrates on the estimation of the lower limb orientation in the sagittal plane. In [26], a combination of inertial sensors and visual data is restricted to the tracking of a single limb (the arm). Moreover, only a simple red arm band is used as image feature. In [24], data obtained from few accelerometers is used to retrieve and play back human motions from a database. [17] presents a system to capture full-body motion using only inertial and

magnetic sensors. While the system in [17] is very appealing because it does not require cameras for tracking, the subject has to wear a suit with at least 17 inertial sensors, which might hamper the movement of the subject. In addition, long preparation time before recording is needed. Moreover, inertial sensors suffer from severe drift problems and cannot provide accurate position information in continuous operation.

1.1. Contributions

Even using learned priors from MoCap data, obtaining limb orientations from video is a difficult problem. Intuitively, because of the cylindrical shape of human limbs, different limb orientations project to very similar silhouettes in the images. These orientation ambiguities can be easily captured by the inertial sensors but accurate positions cannot be obtained. Therefore, we propose to use a small number of sensors (we use only five) fixed at the body extremities (neck, wrists and ankles) as a complementary data source to visual information. On the one hand, we obtain stable and drift-free accurate position information from video data and, on the other hand, we obtain accurate limb orientations from the inertial sensors. In this work, we present how to integrate orientation data from sensors in a contour-based video motion capture algorithm. In several experiments, we show the improved performance of tracking with additional small number of sensors.

2. Twists and Exponential Maps

This section recalls the basics of twists and exponential maps, for further details see [16]. Every 3D rigid motion can be represented by a homogeneous matrix $M \in SE(3)$.

$$M = \begin{pmatrix} R & r \\ 0_{1 \times 3} & 1 \end{pmatrix}, \quad (1)$$

where $R \in SO(3)$ is a rotation matrix and $r \in \mathbb{R}^3$ is a translation. For each matrix $M \in SE(3)$ there is a corresponding twist in the tangent space $se(3)$. An element of $se(3)$ can either be represented by $\theta\xi$, $\theta \in \mathbb{R}$ and $\xi \in \mathbb{R}^6 = \{(v, \omega) | v \in \mathbb{R}^3, \omega \in \mathbb{R}^3, \|\omega\|_2 = 1\}$ or by

$$\theta\hat{\xi} = \theta \begin{pmatrix} \hat{\omega} & v \\ 0_{1 \times 3} & 0 \end{pmatrix} \in \mathbb{R}^{4 \times 4}, \quad (2)$$

where $\hat{\omega}$ is the skew-matrix representation of ω . In the form $\theta\xi$ or $\theta\hat{\xi}$, ξ and $\hat{\xi}$ are referred to as normalized twists, and θ expresses the *velocity* of the twist.

2.1. From Twist to Homogeneous Matrix

Elements from $se(3)$ are mapped to $SE(3)$ using the *exponential map* for twists

$$M = \exp(\theta\hat{\xi}) = \begin{pmatrix} \exp(\theta\hat{\omega}) & (I - \exp(\theta\hat{\omega}))(\hat{\omega}v + \omega\omega^T v\theta) \\ 0_{1 \times 3} & 1 \end{pmatrix} \quad (3)$$

where $\exp(\theta\hat{\omega})$ is the exponential map from $so(3)$ to $SO(3)$ which can be calculated using the Rodriguez formula

$$\exp(\theta\hat{\omega}) = I + \hat{\omega} \sin(\theta) + \hat{\omega}^2 (1 - \cos(\theta)). \quad (4)$$

Note that only sine and cosine functions of real numbers need to be computed.

2.2. Kinematic Chains

The dynamics of the subject are modeled by a *kinematic chain* F , which describes the motion constraints of an articulated rigid body such as the human skeleton [5]. The underlying idea behind a kinematic chain is that the motion of a body segment is given by the motion of the previous body segment in the chain and an angular rotation around a joint axis. Specifically, the kinematic chain is defined with a 6 *DoF* (degree of freedom) root joint representing the global rigid body motion and a set of 1 *DoF* revolute joints describing the angular motion of the limbs. Joints with higher degrees of freedom like hips or shoulders are represented by concatenating two or three 1 *DoF* revolute joints. The root joint is expressed as a twist of the form $\theta\xi$ with the rotation axis orientation, location, and angle as free parameters. Revolute joints are expressed as special twists with no pitch of the form $\theta_j\xi_j$ with known ξ_j (the location and orientation of the rotation axis as part of the model representation). Therefore, the full configuration of the kinematic chain is completely defined by a $(6 + n)$ vector of free parameters

$$\Theta := (\theta\xi, \theta_1, \dots, \theta_n) \quad (5)$$

as described in [18]. Now, for a given point $x \in \mathbb{R}^3$ on the kinematic chain, we define $\mathcal{J}(x) \subseteq \{1, \dots, n\}$ to be the ordered set that encodes the joint transformations influencing x . Let $X = \begin{pmatrix} x \\ 1 \end{pmatrix}$ be the homogeneous coordinate of x and denote π as the associated projection with $\pi(X) = x$. Then, the transformation of a point x using the kinematic chain F and a parameter vector Θ is defined by

$$F_\Theta(x) = \pi(g(\Theta)X) = \pi(\exp(\theta\hat{\xi}) \prod_{j \in \mathcal{J}(x)} \exp(\theta_j\hat{\xi}_j)X). \quad (6)$$

Here, $F_\Theta : \mathbb{R}^3 \rightarrow \mathbb{R}^3$ is the function representing the total rigid body motion $g(\Theta)$ of a certain segment in the chain. Equation (6) is commonly known as the *product of exponentials formula* [16], denoted throughout this paper as F_Θ . In our tracking system, we always seek for differential twist parameters represented in global frame coordinates Θ_d , subsequently we accumulate the motion to obtain the new absolute configuration in body coordinates $\Theta(t)$. Therefore, we have our current configuration at time $t - 1$ given by $\Theta(t - 1)$ and seek for the update Θ_d to find $\Theta(t)$. Recall that $\Theta(t)$ is the vector of twist parameters that represent the map between the body and the global frame at time t . However, at each iteration we update the model and




Figure 2: Absolute twists in continuous line and differential twists in dashed line. X^T denotes a point in global coordinates and X^B denotes a point in body coordinates.

the corresponding twists with the current $\Theta(t-1)$ obtaining the current configuration in global coordinates. Then, we seek for the transformation $g^T(\Theta_d)$ that will transform the model configuration at frame $t-1$ in global coordinates to the model configuration at frame t also in global coordinates. For example, given a point in global coordinates $X^T(t-1)$, we would obtain the point in the next time t as

$$g^T(\Theta_d)X^T(t-1) = X^T(t) \quad (7)$$

where Θ_d are the differential twist parameters at time t in the global frame, see Figure 2. Intuitively, we can think of $g^T(\Theta_d)$ not as a change of coordinates but rather as the twist parameters that give us the instantaneous angular and linear velocity at time t for a point in the global frame. For simplicity, let us denote the differential twist parameters in global coordinates by Θ .

3. Video-based Tracker

In order to relate the surface model to the human's images we find correspondences between the 3D surface vertices and the 2D image contours obtained with background subtraction, see Figure 3. We first collect 2D-2D correspondences by matching the projected surface silhouette with the background subtracted image contour. Thereby, we obtain a collection of 2D-3D correspondences since we know the 3D counterparts of the projected 2D points of the silhouette. In the presented experiments we only use the silhouettes as image features. We then minimize the distance between the transformed 3D points $F_\Theta(X_i)$ and the projection rays defined by the 2D points p_i . This gives us a point-to-line constraint for each correspondence. Defining $L_i = (n_i, m_i)$ as the 3D Plücker line with unit direction n_i and moment m_i of the corresponding 2D point p_i , the point to line distance




Figure 3: (a) Original Image, (b) Background subtracted image, (c) Projected surface mesh after convergence.

d_i can be expressed as

$$d_i = \|F_\Theta(X_i) \times n_i - m_i\| \quad (8)$$

Similar to Bregler *et al.* [5] we now linearize the Equation by using $\exp(\theta\hat{\xi}) = \sum_{k=0}^{\infty} \frac{(\theta\hat{\xi})^k}{k!}$. With I as identity matrix, this results in

$$\pi((I + \sum_{j \in \mathcal{J}(x)} \theta_j \hat{\xi}_j) X_i) \times n_i - m_i = 0. \quad (9)$$

Having N correspondences, we minimize the sum of squared point-to-line distances d_i

$$\arg \min_{\Theta} \sum_{i=1}^N \|d_i\|^2 = \arg \min_{\Theta} \sum_{i=1}^N \|F_\Theta(X_i) \times n_i - m_i\|^2 \quad (10)$$

which after linearization can be re-ordered into an equation of the form $A_1\Theta = b_1$, see Figure 4. Collecting a set of such equations leads to an over-determined system of equations, which can be solved using numerical methods like the Householder algorithm. The Rodriguez formula can be applied to reconstruct the group action g from the estimated twists $\theta_j \xi_j$. Then, the 3D points can be transformed and the process is iterated until convergence. The used video-based tracker is similar to the one presented in [18].

4. Hybrid Tracker

The input of our tracking system consists of:

- Rigid surface mesh of the actor obtained from a laser scanner
- Multi-view images obtained by a set of calibrated and synchronized cameras
- Global orientation data coming from the sensors

We used five inertial sensors fixed at the body extremities (wrists, lower legs, and neck). The final goal is to manipulate the available data in order to relate it linearly (see Figure 4) to the differential kinematic chain parameters Θ that determine the motion from two consecutive frames.




Figure 4: Linear equations derived from orientation data and image silhouettes are combined into a linear equation system.

5. Integration of Sensor Data

5.1. Sensor Data

In our experiments, we use an orientation estimation device MTx provided by XSens [29]. An Xsens MTx unit provides two different streams of data: three dimensional local linear acceleration \ddot{a}^S and local rate of turn or angular velocity ω^S . Orientation data can be obtained from the angular velocity $\omega(k)$ provided by the sensor units. Besides angular velocity, the MTx units provide a proprietary algorithm that can accurately calculate absolute orientations relative to a static global frame F^I , which we will refer to as inertial frame. The inertial frame F^I is computed internally in each of the sensor units in an initial static position and is defined as follows: The Z axis is the negative direction of gravity measured by the internal accelerometer. The X axis is the direction of the magnetic north pole measured by a magnetometer. Finally, the Y axis is defined by the cross product $Z \times X$. For each sensor, the absolute orientation data is provided by a stream of quaternions that define, at every frame, the map or coordinate transformation from the local sensor coordinate system to the global one $q^{IS}(t) : F^S \Rightarrow F^I$. Unfortunately, the world frame defined in our tracking system differs from the global inertial frame. The tracking coordinate frame F^T is defined by a calibration cube placed in the recording volume, in contrast to the inertial coordinate frame which is defined by the gravity and magnetic north directions. Therefore, in order to be able to integrate the orientation data from the inertial sensors into our tracking system, we must know the rotational offset q^{TI} between both worlds, see Figure 5.

Since the Y axis of the cube is perpendicular to the ground and so is gravity, the Y axis of the tracking frame and the Z axis of the inertial frame are aligned. Therefore, q^{TI} is a one parametric planar rotation that can be estimated beforehand using a calibration sequence. Thus, we can easily transform the quaternions so that they define a map from the local sensor frame to the tracking frame F^T :

$$q^{TS} = q^{TI} \circ q^{IS} \quad (11)$$

where \circ denotes quaternion multiplication [20].




Figure 5: Global frames: tracking frame F^T and inertial frame F^I . Local frame: sensor frame F^S .

5.2. Integration of Orientation Data into the Video-based Tracker

In this section we explain how to integrate the orientation data from the sensors as additional equations that can be appended into the big linear system, see Figure 4. Here we have to be very careful and know, at all times, in which frame the rotation matrices are defined. Three coordinate systems are involved: the global tracking frame F^T , the body frame F^B (the local frame of a segment in the chain, e.g. the leg), and the sensor frame F^S . Recall from Sect. 5.1 that the orientation data is given as a quaternion $q^{TS}(t) : F^S \rightarrow F^T$ defining the transformation from the local sensor frame F^S to the global tracking frame F^T , which we will refer to as *ground-truth orientation*. In order to relate the orientation data to the differential twist parameters Θ , we will compare the *ground-truth orientations* $q^{TS}(t)$ of each of the sensors with the estimated sensor orientations from the tracking procedure $\hat{q}^{TS}(t)$, which we will denote as *tracking orientation*. For the sake of simplicity in the operations, we consider from now on the *ground-truth orientation* q^{TS} to be represented as a rotation matrix 3×3 (quaternions can be easily transformed to rotation matrices [20]). The columns of the rotation matrix q^{TS} are simply the sensor basis axes in world coordinates. Let us also define $R(\Theta(t))$ as the total accumulated motion of a body segment at time t , i.e. $R(\Theta(t)) : F^B \rightarrow F^T$. For the sake of clarity we will drop the dependency of Θ and just write $R(t)$. The transformation from the sensor frame to the body frame $q^D(t) : F^S \rightarrow F^B$ is constant during tracking because the sensor and body frame are rigidly attached to the body segment and move together. Thus, we can compute this rotational displacement q^D in the first frame by

$$q^D = R(0)^{-1} q^{TS}(0), \quad (12)$$

where $R(0)$ is the accumulated motion of the body part in the first frame. Now consider the local rotation $R^B(\Theta)$ of frame F^B from time $t - 1$ to time t , see Figure 6. The rotation $R^B(\Theta)$ defined in the body frame is related to the rotation $R^T(\Theta)$ defined in the global frame by the *adjoint*

transformation $Ad_{R^{-1}(t-1)}$

$$R^B(\Theta) = R(t-1)^{-1}R^T(\Theta)R(t-1) \quad (13)$$

Thereby, the *tracking orientation* \hat{q}^{TS} is given by the longer path $F^S \xrightarrow{q^D} F_t^B \xrightarrow{R^B} F_{t-1}^B \xrightarrow{R(t-1)} F^T$, see Figure 6. Now we can compare this transformation matrix to the *ground-truth orientation* given by the sensors q^{TS}

$$R(t-1)R^B(\Theta)q^D = q^{TS}(t). \quad (14)$$

Substituting $R^B(\Theta)$ by its expression in (13) it simplifies to

$$R^T(\Theta)R(t-1)q^D = q^{TS}(t). \quad (15)$$

Therefore, for each sensor s , we can minimize the norm of both matrices with respect to Θ

$$\arg \min_{\Theta} \sum_{s=1}^5 \left\| R_s^T(\Theta)R_s(t-1)q_s^D - q_s^{TS}(t) \right\|. \quad (16)$$

Equation (16) can again be reordered into the form of $A_2\Theta = b_2$ and integrated into the linear system as soft constraints, see Figure 4. Nonetheless, it is interesting to take a closer look at equation (15). Substituting the rotational displacement q^D in equation (15) by its expression in equation (12) we obtain

$$R^T(\Theta)R(t-1)R(0)^{-1}q^{TS}(0) = q^{TS}(t). \quad (17)$$

Expressing $R(t-1)$ in terms of instantaneous rotations

$$R^T(\Theta) \left(\prod_{j=t-1}^0 R^T(j) \right) R(0)^{-1} q^{TS}(0) = q^{TS}(t). \quad (18)$$

Simplifying $R(0)^{-1}$ we obtain

$$R^T(\Theta) \left(\prod_{j=t-1}^1 R^T(j) \right) q^{TS}(0) = q^{TS}(t). \quad (19)$$

This last equation has a very nice interpretation because the columns of the matrix $\left(\prod_{j=t-1}^1 R^T(j) \right) q^{TS}(0)$ are simply the

coordinates of the sensor axis in the first frame (columns of $q^{TS}(0)$), rotated by the accumulated tracking motion from the first frame forward (i.e. not including the initialization motion in frame 0). This last result was very much expected and the interpretation is the following: if we have our rotation matrices defined in a reference frame F^T , we can just take the sensor axes in global coordinates in the first frame (columns of $q^{TS}(0)$) and rotate them at every frame by the instantaneous rotational motions of the tracking. This will result in the estimated sensor axes in world coordinates, which is exactly the *tracking orientation* defined earlier in this Section. Therefore, the problem can be simplified to additional *3D-vector to 3D-vector* constraint equations which can be very conveniently integrated in our




Figure 6: Integration of orientation data into the video-based tracker. *Ground-truth orientation*: clockwise down path from F^S at time t to F^T . *Tracking orientation*: anti-clockwise upper path from F^S at time t to F^T .

twist formulation. Being $\hat{x}(t-1), \hat{y}(t-1), \hat{z}(t-1)$ the *tracking orientation* basis axes in frame $t-1$, and $x(t), y(t), z(t)$ *ground-truth orientation* basis axes in the current frame t , the constraint equations are

$$R^T(\Theta) \begin{bmatrix} \hat{x}(t-1) & \hat{y}(t-1) & \hat{z}(t-1) \end{bmatrix} = \begin{bmatrix} x(t) & y(t) & z(t) \end{bmatrix} \quad (20)$$

which can be linearized similarly as we did in the video-based tracker with image points to mesh points correspondences (*2D-point to 3D-point*). The difference now is that since we rotate vectors, only the rotational component of the twists is needed. For example, the equation for the X-axis correspondence $(\hat{x}(t-1), x(t))$ would be

$$(I + \sum_{j \in \mathcal{J}(x)} \theta_j \hat{\omega}_j) \hat{x}(t-1) = x(t) \quad (21)$$

which depends only on $\theta_j \hat{\omega}_j$. In other words, the constraint equations do not depend at all on the joint axis location nor in the translational motion of the body. This implies that we can integrate the sensor information into the tracking system independently of the initial sensor orientation and location at the body limb.

6. Experiments

In this section, we evaluate our multisensor-fusion approach for motion tracking by comparing the video-based tracker with our proposed hybrid tracker. Learning-based stabilization methods or joint angle limits can also be integrated into the video-based tracker. However, we did not include further constraints into the video-based tracker




Figure 7: Error curves for video-based tracking (red) and hybrid tracking (black), referring to the orientations of the left lower leg for a hopping and jumping motion sequence.

to demonstrate a general weakness of silhouette-based approaches. We note that the video-based tracker works well for many sequences, however in these experiments we focus on the occasions where it fails. Even though benchmarks for video-based tracking are publicly available [22], so far no data set comprising video as well as inertial data exist for free use. Therefore, for our experiments, we generated a data set consisting of 54 takes each having a length of roughly 15 seconds. In total, more than 10 minutes of tracking results were used for our validation study, which amounts to more than 24 thousand frames at a frame rate of 40 Hz. All takes have been recorded in a lab environment using eight calibrated video cameras and five inertial sensors fixed at the two lower legs, the two hands, and the neck. Our evaluation data set comprises various actions including standard motions such as walking, sitting down and standing up as well as fast and complex motions such as jumping, throwing, arm rotations, and cartwheels. For each of the involved four actors, we also generated a 3D mesh model using a laser scanner.

For a given tracking procedure, we introduce a frame-wise *error measure* by considering the angular distance between the two orientations q^{TS} and \hat{q}^{TS} , see Sect. 5.2. This angular distance measured in degrees is defined by the formula

$$d_{quat}(q^{TS}, \hat{q}^{TS}) = \frac{360}{\pi} \arccos \left| \langle q^{TS}, \hat{q}^{TS} \rangle \right|. \quad (22)$$

For a given motion sequence, we compute the error measure for each frame yielding an *error curve*.

In Figure 7, such error curves are shown for two different tracking procedures using the original video-based tracker (red) and the enhanced hybrid tracker (black). For the video-based tracking, there are large deviations between the ground-truth orientations and tracking orientations roughly starting with frame 200. Actually, as a manual inspection revealed, the actor performs in this moment a sudden turn resulting in a failure of the video-based tracking, where the left leg was erroneously twisted by almost 180 degrees. In contrast, the hybrid tracker could successfully track the entire sequence. This is also illustrated by Figure 8. Similarly, the figure also shows a tracking error in the right

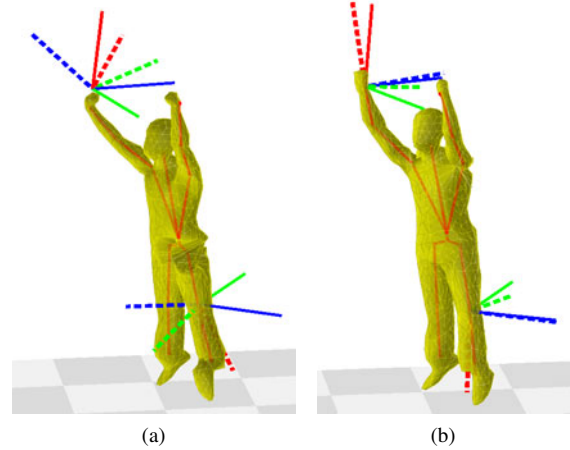


Figure 8: Tracking result for video-based tracking (a) and hybrid tracking (b) for frame 450 of the motion sequence used in Figure 7. Ground-truth orientations in solid lines and tracking orientations in dashed lines.

hand, which is corrected by the hybrid tracker as well. As a second example, we consider a very fast motion, where an actor first rotates his right and afterwards his left arm. Figure 9 shows the error curves for left and right hand for each of the tracking procedures. The curves reveal that the video-based tracker produced significant orientation errors in both hands. This shows that the hand orientations cannot be captured well considering only visual cues. Again, the hybrid tracker considerably improved the tracking results, see also Figure 10. These examples demonstrate how the additional orientation priors resolve ambiguities from image cues. To estimate the quality of our hybrid tracker on more sequences, we computed the error measures (for lower legs, the two hands, and the neck) for each of the five sensors for all sequences and each actor of the data set. A total of 120210 error measures were computed separately for the hybrid and video tracker. We denote mean values and standard deviations of our error measure by μ_V , σ_V and μ_H , σ_H for the video-based and hybrid tracker, respectively. As summarized in Table 1, the sequences of each actor have been improved significantly, dropping the mean error from 30° to 13° . This is also supported by the standard deviations. Let $\tau(s)$ denote the percentage of frames where at least one of the five sensors shows an error of more than s degrees. To show the percentage of corrected severe tracking errors, we computed $\tau_V(45)$ and $\tau_H(45)$ for every actor, see Tab. 1. As it turns out, most of the tracking errors are corrected, dropping the percentage of erroneously tracked frames from 19.29% to 2.51% of all frames. These findings are supported by the normalized histograms of the occurring values of the error measure, see Fig. 11. Furthermore, the hybrid tracker does not increase the computation time




Figure 9: Error curves for video-based tracking (red) and hybrid tracking (black) obtained for an arm rotation sequence (first performed by the right and then by the left arm). **Top:** Left hand. **Bottom:** Right hand.




Figure 10: Tracking result and orientations for two selected frames of the sequence used in Figure 9. (a),(c) Video-based tracking. (b),(d) Hybrid tracking.

of the video-based tracker which is less than 4 s per frame.

One reason for the large amount of corrected errors is that the orientation of limbs is hard to estimate from silhouettes, since the cylindrical shape projects to the same silhouettes in many orientations. By combining the visual with orientation cues, these ambiguities are resolved, resulting in a largely improved performance with the hybrid tracker.

7. Conclusions

In this paper, we presented an approach for stabilizing full-body markerless human motion capturing using a small number of additional inertial sensors. Generally, the goal of reconstructing a 3D pose from 2D video data suffers from inherent ambiguities. We showed that a hybrid approach combining information of multiple sensor types can resolve such ambiguities, significantly improving the tracking quality. In particular, our orientation-based approach could correct tracking errors arising from rotationally symmetric limbs. Using only a small number of inertial sensors fixed at outer extremities stabilized the tracking for the entire underlying kinematic chain.

In the future, we plan to extend our tracker to also make use of acceleration data and rate of turn data, which seem




Figure 11: Normalized histogram, for each actor, of quaternion distances comparison for the whole database.

	Actor 1	Actor 2	Actor 3	Actor 4	Average
μ_V [deg]	26.10	40.80	26.20	31.10	30.29
μ_H [deg]	11.50	14.86	13.98	13.85	13.47
σ_V [deg]	33.79	46.99	29.23	38.07	37.08
σ_H [deg]	9.89	13.01	12.25	14.43	12.28
$\tau_V(45)$ [%]	14.27	29.50	16.53	19.42	19.29
$\tau_H(45)$ [%]	0.47	3.33	2.12	6.45	2.51

Table 1: Mean values μ and standard deviations σ for video-based (V) and hybrid (H) tracker for all sequences of the database, separated by actor. Percentage of large tracking errors denoted by $\tau(45)$.

to be ideally suited to stabilize tracking in outdoor settings, for fast motions, and in the presence of occlusions. To this end, we need suitable strategies that do not destabilize the tracking process in the presence of sensor noise and local artifacts. Furthermore, we want to investigate in how far such fusion techniques make monocular tracking feasible. Finally, we make the multimodal data set used in this paper publicly available at [15] to further support this line of research.

Acknowledgments. This work has been supported by the German Research Foundation (DFG CL 64/5-1 and DFG MU 2686/3-1). Meinard Müller is funded by the Cluster of Excellence on Multimodal Computing and Interaction.

References

- [1] A. Agarwal and B. Triggs. Recovering 3D human pose from monocular images. *IEEE TPAMI*, 28(1):44–58, 2006.
- [2] A. Baak, B. Rosenhahn, M. Müller, and H.-P. Seidel. Stabilizing motion tracking using retrieved motion priors. In *IEEE ICCV*, pages 1428–1435, sep 2009.
- [3] A. Balan, M. Black, H. Haussecker, and L. Sigal. Shining a light on human pose: On shadows, shading and the estimation of pose and shape. In *IEEE ICCV*, volume 1, 2007.
- [4] A. Balan, L. Sigal, M. Black, J. Davis, and H. Haussecker. Detailed human shape and pose from images. In *IEEE CVPR*, pages 1–8, 2007.




Figure 12: Examples of tracking results with our proposed hybrid tracker

- [5] C. Bregler, J. Malik, and K. Pullen. Twist based acquisition and tracking of animal and human kinematics. *IJCV*, 56(3):179–194, 2004.
- [6] D. Brubaker M. A., Fleet and A. Hertzmann. Physics-based person tracking using the Anthropomorphic Walker. In *IJCV (in press)*, 2010.
- [7] E. de Aguiar, C. Stoll, C. Theobalt, N. Ahmed, H. Seidel, and S. Thrun. Performance capture from sparse multi-view video. *ACM Transactions on Graphics*, 27(3):98, 2008.
- [8] H. Dejnabadi, B. Jolles, E. Casanova, P. Fua, and K. Aminian. Estimation and visualization of sagittal kinematics of lower limbs orientation using body-fixed sensors. *IEEE TBME*, 53(7):1382–1393, 2006.
- [9] A. Elgammal and C. Lee. Inferring 3D body pose from silhouettes using activity manifold learning. In *IEEE CVPR*, volume 2, 2004.
- [10] J. Gall, C. Stoll, E. de Aguiar, C. Theobalt, B. Rosenhahn, and H. Seidel. Motion capture using joint skeleton tracking and surface estimation. *IEEE CVPR*, 2009.
- [11] P. Guan, A. Weiss, A. Balan, and M. Black. Estimating Human Shape and Pose from a Single Image. In *IEEE ICCV*, volume 1, 2009.
- [12] L. Herda, R. Urtasun, and P. Fua. Implicit surface joint limits to constrain video-based motion capture. In *LNCS*, volume 3022, pages 405–418, 2004.
- [13] S. Ioffe and D. Forsyth. Human tracking with mixtures of trees. In *IEEE ICCV*, volume 1, pages 690–695, 2001.
- [14] T. Moeslund and E. Granum. A survey of computer vision based human motion capture. *CVIU*, 81(3), 2001.
- [15] Multimodal Human Motion Database MPI08. <http://www.tnt.uni-hannover.de/project/MPI08.Database/>.
- [16] R. Murray, Z. Li, and S. Sastry. *Mathematical Introduction to Robotic Manipulation*. CRC Press, Baton Rouge, 1994.
- [17] D. Roetenberg, H. Luinge, and P. Slycke. Xsens MVN: Full 6DOF Human Motion Tracking Using Miniature Inertial Sensors.
- [18] B. Rosenhahn, T. Brox, and H. Seidel. Scaled motion dynamics for markerless motion capture. In *IEEE CVPR*, pages 1–8, 2007.
- [19] G. Shakhnarovich, P. Viola, and T. Darrell. Fast pose estimation with parameter-sensitive hashing. In *IEEE ICCV*, page 750, 2003.
- [20] K. Shoemake. Animating rotation with quaternion curves. *ACM SIGGRAPH*, 19(3):245–254, 1985.
- [21] H. Sidenbladh and M. Black. Learning the statistics of people in images and video. *IJCV*, 54(1):183–209, 2003.
- [22] L. Sigal and M. Black. HumanEva: Synchronized video and motion capture dataset for evaluation of articulated human motion. Technical Report CS-06-08, Brown University, USA, 2006. Available at <http://vision.cs.brown.edu/humaneva/>.
- [23] L. Sigal, M. Vondrak, and O. Jenkins. Physical Simulation for Probabilistic Motion Tracking. In *IEEE CVPR*, 2008.
- [24] R. Slyper and J. Hodgins. Action capture with accelerometers. In *ACM SIGGRAPH/Eurographics, SCA*, 2008.
- [25] C. Sminchisescu, A. Kanaujia, Z. Li, and D. Metaxas. Discriminative density propagation for 3d human motion estimation. In *IEEE CVPR*, volume 1, page 390, 2005.
- [26] Y. Tao, H. Hu, and H. Zhou. Integration of vision and inertial sensors for 3d arm motion tracking in home-based rehabilitation. *IJRR*, 26(6):607, 2007.
- [27] R. Urtasun, D. Fleet, and P. Fua. 3D people tracking with Gaussian process dynamical models. In *IEEE CVPR*, pages 238–245, 2006.
- [28] J. Wang, D. Fleet, and A. Hertzmann. Gaussian process dynamical models for human motion. *IEEE TPAMI*, 30(2):283–298, 2008.
- [29] Xsens Motion Technologies. <http://www.xsens.com/>, Accessed November 19th, 2009.